

```
In [18]: all <- read.table("hw9t3v0.txt")
colnames(all) <- c('y', 'x1', 'x2', 'x3', 'x4', 'x5')
all[1:10,]
```

y	x1	x2	x3	x4	x5
1.894	1.059	5.074	0.172	2.134	44.277
1.996	1.123	15.074	0.188	2.389	63.779
1.805	2.016	7.695	0.203	3.337	63.008
1.998	1.390	1.316	0.133	3.255	42.985
1.996	1.111	12.508	0.180	2.464	58.507
1.768	1.300	10.594	0.133	2.208	56.031
2.000	1.299	8.195	0.117	2.910	52.984
1.856	1.750	15.555	0.148	2.645	71.427
1.593	1.122	2.348	0.195	3.916	41.664
1.336	1.376	0.961	0.141	3.449	40.332

```
In [31]: n <- 400
q <- 100
train <- all[1:n,]
test <- all[400:500,]
```

После считывания данных и приведения их более-менее приемлимый вид, проведём отбор признаков

```
In [20]: model <- lm(y ~ x1 + x2 + x3 + x4 + x5, data = train)
summary(model)
```

Call:

```
lm(formula = y ~ x1 + x2 + x3 + x4 + x5, data = train)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-1.4923 -0.1956  0.1431  0.2785  0.6240
```

Coefficients:

```
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.628636   0.123226  13.217  < 2e-16 ***
x1           -0.367433   0.075554  -4.863 1.67e-06 ***
x2           -0.021002   0.007408  -2.835 0.00482 **
x3            0.648054   0.587072   1.104 0.27032
x4           -0.066484   0.020512  -3.241 0.00129 **
x5            0.014326   0.003240   4.422 1.27e-05 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.3962 on 394 degrees of freedom

Multiple R-squared: 0.07473, Adjusted R-squared: 0.06299

F-statistic: 6.364 on 5 and 394 DF, p-value: 1.065e-05

Третий признак не является значимым для регрессии, исключим его из модели

```
In [21]: perfect_model <- lm(y ~ model$coefficients[2]x1 + x2 + x4 + x5, data = train)
summary(perfect_model)
```

Call:

```
lm(formula = y ~ x1 + x2 + x4 + x5, data = train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.4724	-0.2133	0.1414	0.2821	0.6241

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.730785	0.081395	21.264	< 2e-16 ***
x1	-0.375911	0.075183	-5.000	8.64e-07 ***
x2	-0.021512	0.007396	-2.909	0.003833 **
x4	-0.068178	0.020460	-3.332	0.000943 ***
x5	0.014700	0.003223	4.561	6.80e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3964 on 395 degrees of freedom

Multiple R-squared: 0.07187, Adjusted R-squared: 0.06247

F-statistic: 7.646 on 4 and 395 DF, p-value: 6.091e-06

```
In [32]: prediction <- predict(perfect_model, newdata=test)
dev <- prediction - test$y
summary(dev)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	-0.6176000	-0.2539000	-0.1455000	-0.0008259	0.2190000	1.2500000

Остатки маленькие, что показывает, что наша модел выбрана хорошо. посмотрим на них

In [33]: `hist(dev)`

